

LAB 1: AI-ASSISTED CITATION INTEGRITY AUDIT

Generate a Research Paper with AI, Then Audit Its Citations

AI Tools for Research • 2 Hours • Individual Activity • Topic IDs T1-T40

Student Detail	Fill in
Name	Tushar Sukumar Sarkar
Roll Number	PSE2026001
Programme / Department	Information Technology in Software Engineering and Automation
Topic ID	3
Full Assigned Topic	Token-Level Uncertainty as a Predictor of Hallucination in Large Language Models
AI Tool Used	Claude
Date	21/08/2026

IMPORTANT: You will receive a topic from T1-T40. Your task is not to produce the “best” paper. Your task is to generate an AI-written research paper and systematically determine whether its citations can be trusted.

1. PURPOSE OF THIS LAB

Generative AI systems can produce research papers that look scholarly and may contain convincing references, author names, journal titles, DOIs, arXiv IDs, and in-text citations. However, the presence of a citation does not guarantee that the cited publication actually exists, that its metadata are correct, that the identifier belongs to that publication, or that the cited paper supports the statement made by the AI.

In this laboratory activity, you will test these issues yourself by generating a research paper and auditing a systematic sample of its references.

Question 1 - Citation Authenticity: Does the cited publication actually exist, and is it correctly identified?

Question 2 - Claim-Citation Support: If the publication is genuine, does it actually support the claim for which the AI cited it?

2. IMPORTANT RULES

- Use only the topic assigned to you by the instructor.
- Generate the paper in a fresh AI conversation.
- Do not tell the AI that you are conducting a citation audit or hallucination test.
- Do not intentionally ask the AI to generate fake references.
- Preserve the original AI-generated paper exactly as produced.
- Do not correct, regenerate, replace, or question references before saving the original output.
- For verification, do not accept another AI chatbot's statement as proof that a reference exists.
- Use scholarly databases, publisher pages, DOI/arXiv/PMID records, or other reliable scholarly sources.
- Record what you actually find, even if all references are genuine or all are problematic.

Most important principle: You are not being evaluated on whether the AI produces good or bad references. You are being evaluated on how systematically and correctly you audit them.

3. WHAT YOU WILL DO

Part	Activity
------	----------

A	Record your baseline perception of AI citations
B	Record the AI environment used
C	Generate and preserve an AI-written research paper
D	Create a reference inventory
E	Select 10 references systematically
F	Judge their plausibility before verification
G	Verify authenticity, metadata and identifiers
H	Determine whether genuine references support the AI-generated claims

4. RECOMMENDED 2HOUR TIME PLAN

Time	Activity
0-10 min	Parts A-B: Baseline and AI environment
10-30 min	Part C: Generate and save research paper
30-40 min	Parts D-E: Inventory and select references
40-45 min	Part F: Pre-verification plausibility
45-80 min	Part G: Authenticity and metadata audit
80-110 min	Part H: Claim-citation entailment audit
110-120 min	Calculate scores, check worksheet and prepare submission

Do not spend excessive time formatting the AI-generated paper. The audit is the main activity.

PART A- BASELINE ASSESSMENT

Complete this section BEFORE opening the AI tool.

A1. How often do you use generative AI for academic work?

☐ Never ☐ Rarely ☐ Occasionally ☒ Weekly ☐ Almost daily

A2. Have you previously asked an AI system to generate references?

☐ Frequently ☒ Occasionally ☐ Once or twice ☐ Never

A3. Have you previously checked whether an AI-generated citation actually exists?

☐ Always ☐ Sometimes ☐ Rarely ☒ Never ☐ Not applicable

A4. Before this lab, how confident are you that citations generated by AI are generally reliable?

☐ 1 ☒ 2 ☐ 3 ☐ 4 ☐ 5 (1 = not at all confident; 5 = completely confident)

Important: Do not change this answer later, even if your opinion changes during the experiment.

PART B- RECORD YOUR AI ENVIRONMENT

Information	Your Response
AI tool used	Claude
Model name	Haiku
Model/version shown	4.5 Extended
Access tier	☒ <u>Free</u> ☐ Paid/Pro ☐ Institutional ☐ Not sure
Web browsing available?	☒ <u>Yes</u> ☐ No ☐ Not sure
Was web browsing actually used?	☒ <u>Yes</u> ☐ No ☐ Not sure
Research/Deep Research mode used?	☐ Yes ☒ <u>No</u>
Date/time generated	

Important: If the model/version is not displayed, write NOT SHOWN. Do not guess the model version.

PART C- GENERATE THE AI RESEARCH PAPER

C1. Open a Fresh AI Conversation

Do not tell the AI that this is a citation audit, that you are testing hallucinations, or that you expect fabricated references. We want to observe normal AI research-paper generation behaviour.

C2. Use the Following Common Prompt

Replace only [YOUR ASSIGNED TOPIC].

Generate a scholarly research paper on the topic: “[YOUR ASSIGNED TOPIC]”. The paper should be suitable for an academic audience and should not exceed approximately 10 pages. Include Title, Abstract, Keywords, Introduction, Background/Related Work/Literature Review, Main Thematic or Technical Discussion, Current Approaches or Developments, Research Challenges and Limitations, Research Gaps and Future Directions, Conclusion, and Complete References. Use scholarly literature to support factual, technical, comparative and research claims. Provide in-text citations throughout the paper. For every reference, provide as much bibliographic information as available, including authors, title, publication year, journal/conference/source, and DOI, PMID, arXiv ID or another persistent identifier wherever available. Use genuine scholarly sources. Do not intentionally fabricate references.

C3. Preserve the Original Output

☐ Save the complete AI-generated paper

☐ Save the complete reference list

The screenshot shows a document editor window titled "Token_Level_Uncertainty_Hallucination_Research_Paper.docx". The document content includes:

However, significant challenges remain before uncertainty-based hallucination mitigation becomes standard practice. Calibration brittleness across domains, computational efficiency requirements, heterogeneous hallucination types, and incomplete theoretical understanding all present obstacles to widespread deployment. The field lacks unified evaluation frameworks and systematic investigation of the mechanisms linking uncertainty to hallucination.

The most promising research direction likely involves integrated approaches combining multiple signals: token-level uncertainty, semantic coherence, knowledge consistency checks, and retrieval augmentation. Rather than relying solely on confidence scores, robust hallucination mitigation will likely require ensemble strategies leveraging complementary information sources.

For practitioners, current uncertainty-based approaches show sufficient promise to warrant experimentation, particularly semantic uncertainty methods for domain-specific applications where calibration can be validated. However, reliance on uncertainty as a silver-bullet hallucination solution is premature. The field should continue developing the theoretical understanding, evaluation methodologies, and practical techniques necessary to realize the potential of uncertainty-based approaches while managing their fundamental limitations.

8. References

1. Brown, T., Mann, B., Ryder, N., et al. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33, 1877-1901.
2. Geifman, Y., & El-Yaniv, R. (2017). Selective prediction using unreliable classifiers. *Advances in Neural Information Processing Systems*, 30, 3754-3764.
3. Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2016). On calibration of modern neural networks. *International Conference on Machine Learning (ICML)*, 70, 1321-1330. arXiv:1706.04599
4. Huang, L., Yu, W., Ma, W., et al. (2023). A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. arXiv preprint arXiv:2311.05232.
5. Kadavath, S., Conerly, T., Asks, T., et al. (2022). Language models (mostly) know what they know. arXiv preprint arXiv:2207.05221.
6. Kendall, A., & Gal, Y. (2017). What uncertainties do we need in Bayesian deep learning for computer vision? *Advances in Neural Information Processing Systems*, 30, 5574-5584. arXiv:1703.04977
7. Kuhn, L., Gal, Y., & Farquhar, S. (2023). Semantic uncertainty: Linguistic invariances for uncertainty estimation in natural language generation. arXiv preprint arXiv:2302.09664.
8. Lakshminarayanan, B., Pritzel, A., & Blundell, C. (2017). Simple and scalable predictive uncertainty estimation using deep ensembles. *Advances in Neural Information Processing Systems*, 30, 6402-6413. arXiv:1706.04599
9. Lewis, P., Perez, E., Rinott, R., et al. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9457-9474. arXiv:2005.11401
10. Malkin, N., Yildirim, I., & Salakhutdinov, R. (2023). Measuring factuality in generation with dependency-level entailment. *Findings of the Association for Computational Linguistics: EMNLP 2023*, 14487-14507. arXiv:2309.09556
11. Maynez, J., Narayan, S., Artetxe, M., & Schwenk, H. (2021). On faithfulness and factuality in abstractive summarization. *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)*, 1906-1919. arXiv:2104.13346
12. Petroni, F., Rocktäschel, T., Lewis, P., et al. (2019). Language models as knowledge bases? *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2463-2473. arXiv:1909.01066
13. Roller, S., Suhr, A., Lin, S., et al. (2023). Evaluating language models for indexing scholarly citations. arXiv preprint arXiv:2303.07302.
14. Schuster, T., Fisch, A., Barzilay, R., & Adar, E. (2022). Confident natural language inference. *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2869-2880. arXiv:2210.03387
15. Thawani, A., Wangni, J., & Raffel, C. (2022). Questioning the AI: Informing design practices for explainable AI user experiences. *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems*, 1-17.
16. Touvron, H., Martin, L., Stone, K., et al. (2023). Llama 2: Open foundation and fine-tuned chat models. arXiv preprint arXiv:2307.09288.
17. Yildirim, I., Sap, M., LeBras, R., et al. (2021). Do language models have coherent mental models of physical events? *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 4988-5005. arXiv:2011.13762
18. Zhang, Y., Marone, M., Li, Y., et al. (2023). Evaluating hallucinations in language models. arXiv preprint arXiv:2311.15930.
19. Zhong, M., Liu, P., Wang, Y., & Huang, M. (2020). Extractive summarization as text matching. *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 6197-6209. arXiv:2004.08795

The screenshot also shows a Windows taskbar at the bottom with the time 22:48 and date 21-08-2026.

☐ Take screenshot(s) showing the generated output/references

<https://claude.ai/share/4cd2c8c6-bc93-40da-a7ee-6f22b0910fec>

☐ Save conversation/export/share evidence if available

STOP: Do not ask the AI to correct or verify anything yet. Your original output is the experimental evidence.

PART D- PAPER PROFILE AND REFERENCE INVENTORY

Paper title:

Measure	Value
Approximate pages	8
Approximate word count	3000
Number of sections	9 including Reference
Total number of references	19
Approximate in-text citation occurrences	
Did AI warn you to verify references?	☐ Clearly ☐ Vaguely ☒ No

D1. Reference Inventory

Classify every reference once using the best available identifier in this priority order: DOI -> PMID -> arXiv ID -> URL only -> No persistent identifier. The category totals must equal the total number of references in the AI-generated paper.

Reference Type	Count
DOI available	0
PMID available, no DOI	0
arXiv ID available, no DOI/PMID	16
URL only	0
No persistent identifier supplied	3
TOTAL	19

Examples

Category	Example	How to Count
DOI available	Kather, J. N., et al. (2019). Predicting survival from colorectal cancer histology slides using deep learning. PLoS Medicine. DOI: 10.1371/journal.pmed.1002730	DOI available
PMID available, no DOI	A biomedical paper for which the AI provides PMID: 12345678, but no DOI is supplied/found	PMID available, no DOI
arXiv ID available	Yao, S., et al. ReAct: Synergizing Reasoning and Acting in Language Models. arXiv:2210.03629	arXiv ID available
URL only	World Health Organization. Colorectal cancer. Available at: https://... with no DOI, PMID or arXiv ID	URL only
No persistent identifier	Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep Learning. MIT Press.	No persistent identifier if none is supplied

Important: Absence of a DOI does not mean a reference is fake. An identifier being present also does not prove authenticity; it must be checked.

PART E- SELECT 10 REFERENCES FOR AUDIT

You are not allowed to choose only suspicious-looking references. This prevents selection bias.

If the bibliography contains 10 or fewer references: audit all references.

If it contains more than 10 references: select the first 3 references, the last 3 references, and 4 references approximately evenly distributed through the middle.

Example: If there are 28 references, a reasonable sample could be 1, 2, 3, 8, 13, 18, 23, 26, 27, 28.

Audit Slot	Original Bibliography Reference No.
R1	1
R2	2
R3	3
R4	6
R5	9
R6	12
R7	15
R8	17
R9	18
R10	19

Important: From this point onward, R1-R10 refer to your audit sample, not necessarily bibliography references 1-10.

PART F- PLAUSIBILITY TEST

Do this BEFORE searching for any reference.

Look only at the citation as generated by the AI. Ask yourself: “If I had not verified this citation, how convincing would it look to me?”

Score	Meaning
1	Obviously suspicious
2	Somewhat suspicious
3	Unsure
4	Looks genuine
5	Completely convincing

Ref.	Initial Impression	Plausibility 1-5	Predict Genuine?
R1	Seems Suspicious	2	No
R2	suspicious	2	No
R3	genuine	4	Yes
R4	Not sure	3	Yes
R5	Looks Genuine	5	Yes
R6	genuine	4	Yes
R7	genuine	4	Yes
R8	Not sure	3	No
R9	genuine	5	Yes
R10	Seems genuine	4	Yes

Example: A citation with realistic authors, a well-known journal, complete volume/pages and a DOI-like identifier may look highly convincing. You might record Plausibility = 5 and Predict Genuine = Yes. You are recording your pre-verification belief, not the truth.

Critical rule: Never change Part F after completing Part G. Your incorrect predictions are valuable experimental observations.

PART G- CITATION AUTHENTICITY AND METADATA AUDIT

G1. What Are You Checking?

- Does the publication exist?
- Is the title correct?
- Are the authors correct?
- Is the year correct?
- Is the journal/conference/source correct?
- Are volume, issue and pages correct, if supplied?
- Does the DOI/arXiv/PMID belong to this publication?

G2. Verification Procedure

☐ Persistent identifier: DOI / arXiv / PMID

☐ Publisher or official repository

☐ Google Scholar

☐ Semantic Scholar / OpenAlex

☐ Crossref

☐ PubMed or another discipline-specific database

☐ Exact-title search

☐ Author + title-keyword search

Not acceptable as sole evidence: “I asked another AI and it said the paper exists.” AI can help you search, but final verification must rely on scholarly or publisher evidence.

G3. Assign One A-E Code

Code	Classification	Meaning
A	VERIFIED	Publication exists and important metadata are correct.
B	WRONG METADATA	Publication exists, but one or more bibliographic details are incorrect.
C	FRANKENSTEIN	Real-looking pieces are combined, but that specific publication does not exist.
D	FABRICATED	No reliable scholarly evidence that the publication exists.
E	IDENTIFIER MISMATCH	DOI/arXiv/PMID is invalid or resolves to a different publication.

G4. Examples

A: AI gives a paper and all details match the publisher record.

B: The paper exists, but the AI gives the wrong publication year or journal.

C: The authors are real, the journal is real, and the topic is plausible, but no such paper with that title/authorship combination exists.

D: You cannot find the paper or a credible scholarly record after systematic searching.

E: The DOI supplied by AI opens an entirely different paper or does not resolve.

G5. Record Your Evidence

Ref.	Publication Found?	Title Match?	Authors Match?	Year Match?	Venue Match?	Identifier Match?
------	--------------------	--------------	----------------	-------------	--------------	-------------------

R1	Y	Y	Y	Y	NA	NA
R2	Y	Y	Y	Y	NA	NA
R3	Y	Y	Y	N	NA	Y
R4	Y	Y	Y	Y	NA	Y
R5	Y	Y	Y	Y	NA	Y
R6	Y	Y	Y	Y	NA	Y
R7	Y	Y	N	N	NA	NA
R8	Y	Y	Y	N	NA	N
R9	Y	Y	Y	N	NA	N
R10	Y	Y	Y	Y	NA	Y

Ref.	Verification Source(s)	Identifier Checked	What You Found
R1	[2005.14165] Language Models are Few-Shot Learners	[PDF] Language Models are Few-Shot Learners Semantic Scholar	Introduces GPT-3 and establishes the foundational capabilities of massive auto-regressive language models to perform tasks via few-shot prompting.
R2	[1705.08500] Selective Classification for Deep Neural Networks	(PDF) Selective Classification for Deep Neural Networks	Discusses selective prediction (allowing a model to abstain from answering when uncertain), which is the theoretical basis for early-exit strategies in hallucination mitigation.
R3	[1706.04599] On Calibration of Modern Neural Networks	[PDF] On Calibration of Modern Neural Networks Semantic Scholar	Analyzes the "calibration problem" in deep learning, explaining why modern neural networks are often overly confident even when generating incorrect outputs.
R4	[1703.04977] What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision?	[PDF] What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision? Semantic Scholar	Provides the critical mathematical distinction between data inherent and model ignorance uncertainty, which is essential for diagnosing why an LLM hallucinates.
R5	[2005.11401] Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks	[PDF] Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks Semantic Scholar	Introduces Retrieval-Augmented Generation (RAG), a key architecture where external context is provided to ground the model and reduce factual fabrications.
R6	[1909.01066] Language Models as Knowledge Bases?	[PDF] Language Models as Knowledge Bases? Semantic Scholar	Explores probing language models to see if they function as relational knowledge bases, setting the groundwork for evaluating what factual data a model actually "knows" versus what it fabricates.
R7	Questioning the AI: Informing Design Practices for Explainable AI User Experiences Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems	Questioning the AI: Informing Design Practices for Explainable AI User Experiences	Focuses on Explainable AI (XAI) and user design, highlighting how an AI's internal uncertainty and potential for error should be communicated to human users.

R8	[2212.10029] Do language models have coherent mental models of everyday things?	2212.10029	Investigates whether language models possess coherent internal "world models," addressing a root cause of hallucinations: the lack of grounded physical or logical reasoning.
R9	[2310.03368] Evaluating Hallucinations in Chinese Large Language Models	A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions	Provides a contemporary, direct evaluation framework specifically targeted at defining, categorizing, and measuring hallucinations in modern LLMs.
R10	[2004.08795] Extractive Summarization as Text Matching	[PDF] Extractive Summarization as Text Matching Semantic Scholar	Frames text summarization as a semantic matching problem, providing methodologies useful for clustering generated responses and measuring semantic entropy (checking if varied outputs mean the same thing).

G6. Final Classification

Ref.	Final Code	One-Line Reason
R1	A	The paper exists, and all details perfectly match the official NeurIPS 33 publisher record.
R2	C	The authors published a highly relevant paper in NeurIPS 2017, but its actual title is "Selective Classification for Deep Neural Networks"; the generated title is a plausible fabrication
R3	B	The paper exists and the authors/title are correct, but it was published and presented at ICML in 2017, not 2016.
R4	A	The paper exists, and all details perfectly match the official NeurIPS 30 publisher record.
R5	A	The paper exists, and all details perfectly match the official NeurIPS 33 publisher record.
R6	A	The paper exists, and all details perfectly match the official EMNLP 2019 conference record.
R7	C	The title belongs to a well-known 2020 CHI paper authored by Q. Vera Liao et al.; the provided authorship combination is completely fabricated.
R8	E	The provided identifier (ar Xiv: 2011. 13762) resolves to an entirely different paper about mathematics ("Fourier duality in the Brascamp-Lieb inequality").
R9	E	he provided identifier (ar Xiv: 2311. 15930) resolves to an entirely different paper ("WorldSense: A

		Synthetic Benchmark for Grounded Reasoning in Large Language Models").
R10	A	The paper exists, and all details perfectly match the official EMNLP 2020 conference record

Good reason: B: Article exists, but AI reported 2022 instead of 2020.

Poor reason: B: Details wrong.

PART G- SCORING

Count your final classifications: A = 5 B = 1 C = 2 D = 0 E = 2

Total audited: N = A + B + C + D + E = 10

Code	Points
A	4
B	3
C	1
D	0
E	1

Authenticity Points = $4A + 3B + C + E$

Authenticity Score = $[(4A + 3B + C + E) / (4N)] \times 100$

My calculation: $4(5) + 3(1) + 1(2) + 0(0) + 1(2) = 27$ points

Authenticity Score = $27 / (4 \times 10) \times 100 = 67.5\%$

Worked example: If A=6, B=2, C=1, D=1, E=0, then points = $4(6)+3(2)+1(1)=31$; maximum = 40; Authenticity Score = $31/40 \times 100 = 77.5$.

Score	Interpretation
90-100	Excellent authenticity
75-89	Good, minor concerns
60-74	Moderate citation-authenticity risk
40-59	High citation-authenticity risk
Below 40	Very high citation-authenticity risk

G7. Prediction Accuracy

Compare your Part F Genuine Yes/No predictions with the Part G results.

Prediction Accuracy = $(\text{Correct genuine/fake predictions} / \text{Total predictions}) \times 100$

Correct predictions = 7 Total predictions = 10 Prediction Accuracy = 70 %

FINAL LAB 1 RESULTS

Measure	Your Result
Total references in AI paper	19
References deeply audited	10
A: Verified	5

B: Wrong metadata	1
C: Frankenstein	2
D: Fabricated	0
E: Identifier mismatch	2
Authenticity Score	___67.5 /100
Prediction Accuracy	___70 %

FINAL INTERPRETATION

1. Did professional-looking citations necessarily turn out to be genuine?

Not necessarily few of them were genuine and few of them were fake

2. Did all genuine references actually support the claims for which they were cited?

In my case if we remove the fake ones other were supporting the claim for which it were cited

3. What was the most serious citation-integrity problem you observed?

Author was totally mismatched with the paper and identifier pointed to a totally different paper

4. What is the single most important lesson you learned from this lab?

Not to trust blindly in ai and cross check everything if using it, be more responsible with it

SUBMISSION INSTRUCTIONS

Submit one PDF: Lab1_RollNumber_Name


Example: Lab1_IEC2026001_RahulSharma

STUDENT DECLARATION

I confirm that I generated the research paper using the recorded AI environment, preserved the original output before conducting the audit, followed the prescribed reference-selection procedure, and personally verified the evidence reported in this worksheet. I have not altered my pre-verification predictions after seeing the verification results.

Student Name:Tushar Sukumar Sarkar

Roll Number:PSE2026001



Signature:

Date: 21/08/2026

CENTRAL LESSON OF LAB 1

Does the source exist? + Does the source actually support the claim?

AI Generation -> Reference Inventory -> Systematic Sampling -> Prediction -> Authenticity Verification -> Claim-Support
Verification -> Quantitative Scores -> Submission